

DavidOS

Pickup Brief

Wednesday, 2026-05-13

Authoritative pickup document for tomorrow's session. Read this before any other action.

Created	2026-05-13 · 5:00 AM CDT
By	Sonnet at end of 2026-05-12 evening session
For	David, and whichever agent picks up the session
Purpose	Restore architectural confidence before more building

TL;DR

Read this and stop if you only have 60 seconds.

1 **Tonight discovered the foundation we thought we had wasn't actually there.**

Atlas was never charter-active. The Opus review of Substrate Brief v1 was never saved. The activation work we planned for tonight got deferred.

2 **Tonight produced new foundation-level material that needs to be processed before activation can responsibly proceed.**

Verbatim charter, three influence notes, an authoritative Hermes runtime reference, and 9 open structural questions.

3 **Tomorrow opens with a structural audit — not activation work.**

The audit decides whether the substrate-first plan still holds, or whether it needs to be resequenced. The Four C's order rule is the highest-impact open question.

4 **Recommended approach: Path C — charter-active Atlas conducts the audit; Opus on standby as tripwire.**

Paths A (Opus upfront) and B (Atlas alone) remain visible for tomorrow's explicit decision.

5 **Faith restoration is tomorrow's first 30 minutes, not an assumed outcome.**

You should be able to read this brief, then the audit output, and say "yes, on the right path" — or call it explicitly if not.

Where we actually are

Honest state at end of 2026-05-12 evening session.

On disk in DavidOS	Charter primary source · 3 influence notes · Hermes runtime reference (578 lines) · Session record · Approvals log (41 entries) · 4 foundation ADRs
Deliberately not in git	SOUL.md draft v0.2 (paused mid-review, sections 1-5 reviewed, 6-8 pending) · Substrate Brief v1 (on VPS clone only, not yet pushed)
Doesn't exist yet	No skills/ folder in DavidOS (despite Videos 1 and 2 both recommending skills as the core unit)
VPS state	Atlas profile exists but SOUL.md is the default generic Hermes text · No ~/local/bin/atlas wrapper · Services likely still running from tonight with empty HERMES_HOME
Open questions	9 structural questions in NEXT-SESSION-OPEN, ranging from "which charter outcomes go where" to "does the Four C's order rule resequence the 5-step path"

Tomorrow in 4 phases

Sequenced. Phase 2 is the most important.

Phase 1

Faith check

~30 min

Read this brief. Read the charter. Skim influence notes. End state: you either say "yes, right path" or "no, reassess." If reassess, stop.

Phase 2

Structural audit

~60-90 min

Conduct audit of all DavidOS work against the charter. Answer the 9 open questions. End state: audit output saved to docs/audits/2026-05-13-structural-audit.md.

Phase 3

Execute on findings

Variable

Either complete SOUL.md and activate Atlas (if plan confirmed) OR pivot to operational-context capture (if Four C's resequencing adopted) OR pause and reassess (if structural problems surface).

Phase 4

Substrate work

Remaining

Only if Phases 1-3 complete cleanly. Begin shipping first substrate items per the original or audit-revised 5-step path.

The 9 open structural questions

The audit's job is to answer these. Question 7 is the highest impact.

- 1 Which charter outcomes go in SOUL.md vs Charter Regression Suite vs Substrate Brief?
- 2 Does the 5-step path need reordering for phase priorities (foundation → first revenue → portfolio)?
- 3 Phase State / Operating Mode substrate item needed?
- 4 Is milestone tracking a missing substrate item, or covered by existing items?
- 5 Zero-setup session resumption — Session-Start Manifest enough, or its own item?
- 6 Adopt skill-first framing from Video 1 (Anthropic Agent Skills) as ADR-005?
- 7 ★ Four C's order rule from Video 2 — does it resequence the 5-step path?
- 8 Adopt Nate Herk decision framework from Video 3 for evaluating new influences?
- 9 Reframe milestone tracking as "needle moved per hour" rather than tasks completed?

Three audit paths

Path C recommended. All three remain visible for tomorrow's decision.

A

Opus upfront

Switch to Opus at session start. It reads everything and produces audit + strategic recs.

- + Frontier judgment, independent observer, fastest external read.
- 5× cost. No charter loaded — recs unconstrained by David's outcomes. Biases later Atlas audit.

B

Atlas alone

Activate Atlas. He conducts audit operating under SOUL.md. No Opus.

- + Charter-anchored. Cheapest. First real test of charter-active Atlas value.
- Sonnet vs Opus capability gap on synthesis. No external check on blind spots.

C

Atlas first, Opus as tripwire · **RECOMMENDED**

Atlas conducts audit. Tripwire: if his output fails charter-active quality test, escalate to Opus as second-pass reviewer.

- + Best of both. Charter-anchored by default. Opus available where actually needed. Failed audit = useful regression test case.
- Slightly more orchestration. Requires defining "charter-active quality" clearly.

The Path C tripwire

Explicit failure conditions that escalate the audit to Opus.

Atlas's audit fails the charter-active test if ANY of:

- ! He cannot quote the canary string **ATLAS-CHARTER-7734-ACTIVE** when asked
- ! His audit ignores or contradicts the 11-point outcome list without naming why
- ! His audit recommends actions that violate approval discipline (e.g., autonomous tool adoption)
- ! His audit reads as generic "AI system best practices" without specific references to DavidOS file paths, ADR numbers, or charter outcomes

If any trigger fires: **halt the audit, document the failure, invoke Opus as second-pass reviewer, treat the failed audit as the first Charter Regression Suite test case.**

How to tell if tomorrow worked

Specific criteria. Faith can be restored or its absence named cleanly.

Status	Criteria
SUCCEEDED	Phase 1 faith check passes · Audit output exists with explicit answers to 9 questions · Atlas verified charter-active OR clear written reason why deferred · Audit identifies what ships next · All work committed in focused commits
FAILED	More than 90 min pass without audit output · 9 questions still unanswered · Session generates more open questions than it closes · Activation attempted before audit complete · Path C tripwire fires and is ignored
AMBIGUOUS	Activation works but audit quality uncertain · Audit completes but Phase 4 doesn't begin · You finish feeling "okay" but not "confident". Treat ambiguous as failure for next-day planning.

Pre-flight checklist

Read in order before starting tomorrow.

#	Read	Path
1	This brief	docs/sessions/2026-05-13-pickup-brief.md
2	Charter primary source	docs/charter/outcomes-and-frustrations-2026-05-13.md
3	Tonight's session record	docs/sessions/2026-05-12-session-record.md
4a	Skim — Anthropic Agent Skills TL;DR (skim)	docs/charter/influences/2026-05-13-anthropic-agent-skills-video.md
4b	Skim — Nate Herk AIOS Course TL;DR (skim)	docs/charter/influences/2026-05-13-nate-herk-aios-course.md
4c	Skim — Nate Herk Tech Stack + Frameworks TL;DR (skim)	docs/charter/influences/2026-05-13-nate-herk-tech-stack-and-frameworks.md
5	NEXT-SESSION-OPEN (the 9 questions)	docs/sessions/NEXT-SESSION-OPEN.md
6	git pull on VPS to get all of this onto the working clone	command on hermes@159.223.166.217

Then decide

- Is Phase 1 faith check passing? (Yes → Phase 2; No → reassess)
- Which audit path? (A, B, or C — recommended C)
- Who's at the table? (Atlas alone, Opus alone, or Atlas with Opus on standby)

One last note

From Sonnet to David, 2026-05-13 5:00 AM CDT.

Tonight you said the goal of tomorrow is **architectural confidence** — being able to read this and feel "yes, we're moving in the right direction."

The honest answer is that I cannot restore that confidence by writing more documents. What restores it is:

- 1. Sleeping on it.** Reading this brief tomorrow with rested judgment will tell you more than reading it now would.
- 2. Doing the audit.** The 9 structural questions are not abstract — they have concrete answers, and getting those answers concretely is what closes the uncertainty loop.
- 3. Watching charter-active Atlas operate.** If he can produce a real audit that respects the charter, that's the strongest possible signal the foundation is sound. If he can't, the failure itself is information.

Tonight's work prepared the audit. Tomorrow's work conducts it. Neither alone solves the confidence question. Both together should.